

Topic 1: Overview of the Network Layer (Data Plane vs. Control Plane)

The primary role of the network layer is to move packets from a sending host to a receiving host. Unlike the transport and application layers, network-layer protocols run in each and every host and router within the network. The network layer can be decomposed into two distinct, interacting parts: the **data plane** and the **control plane**.

1. Core Network Layer Functions

To move packets from sender to receiver, the network layer executes two key functions:

- **Forwarding (Data Plane):**
 - Refers to the local, per-router action of transferring a packet from an input link interface to the appropriate output link interface.
 - Occurs at extremely short timescales (typically a few nanoseconds).
 - Is almost entirely implemented within the router's hardware.
 - Relies on a **forwarding table** inside the router, which is indexed using packet header fields to determine the correct outgoing link interface.
- **Routing (Control Plane):**
 - Refers to the network-wide process that determines the end-to-end path or route taken by packets from a source host to a destination host.
 - Occurs at longer timescales (typically seconds).
 - Is often implemented in software.
 - Utilizes **routing algorithms** to calculate these optimal paths.

The Driving Analogy:

- **Forwarding** is like a driver navigating through a single highway interchange (choosing which ramp to exit).
- **Routing** is like planning the entire trip from start to finish using a map before embarking on the journey.

2. Control Plane Approaches: Traditional vs. SDN

The network layer handles the configuration of local forwarding tables using one of two architectural approaches:

Feature	Traditional Approach	Software-Defined Networking (SDN) Approach
Implementation	Monolithic configuration where both routing and forwarding are contained within each individual router.	Explicit separation of the data plane and control plane functions.
Routing Control	A routing algorithm runs in every router. Routers exchange routing messages using a routing protocol to compute forwarding tables.	A physically separate, centralized remote controller computes and distributes the forwarding tables.
Router Execution	The router executes routing algorithms locally in software and forwards packets in hardware.	The physical router performs packet forwarding only, while the control logic is completely handled by the remote software controller.

3. Explanatory Figures

Figure 1.1: Functional Split of Data Plane vs. Control Plane






Figure 1.2: Traditional vs. SDN Architectural Comparison

Traditional Monolithic Router



SDN (Software-Defined Networking)



Would you like to move on to the next topic, which covers the Network Service Model and an introduction to the internal hardware architecture of a router?



Topic 2: Network Service Models & Router Architecture

1. Network Service Models

The network service model defines the exact characteristics of end-to-end packet delivery between a sending host and a receiving host. Depending on the network architecture, the network layer can offer various types of service guarantees:

- **Guaranteed delivery:** Guarantees that every packet sent by a source host will eventually reach the destination host.


- **Guaranteed delivery with bounded delay:** Guarantees that packets are delivered within a strictly specified host-to-host delay bound (e.g., within 100 msec).
- **In-order packet delivery:** Guarantees that packets arrive at the destination in the exact order that they were sent.
- **Guaranteed minimal bandwidth:** Emulates a dedicated transmission link of a specified bit rate (e.g., 1 Mbps) between the hosts. Packets are guaranteed delivery as long as the sending host transmits below this bit rate.
- **Security:** Encrypts all datagrams at the source and decrypts them at the destination to provide end-to-end confidentiality.

The Internet vs. Alternative Service Models

Network Architecture	Service Model	Guarantees Provided
The Internet (IP)	Best-effort service	None. No guarantees on packet delivery, in-order arrival, end-to-end delay, or minimal bandwidth. Despite this lack of guarantees, adequate bandwidth provisioning has proven "good enough" to support streaming video and real-time conferencing apps.
ATM	ATM Service Model	Guarantees in-order delivery, bounded delay, and a minimal bandwidth capacity.
Intserv (Proposed IP Extension)	Integrated Services	Aims to provide absolute end-to-end delay guarantees and congestion-free communication paths.

2. General Router Architecture Overview

A generic network router consists of four primary internal components:



1. **Input Ports:** Microprocessors that handle physical link termination, link-layer protocol interoperation, and local forwarding lookups.
2. **Switching Fabric:** The internal routing components that physically connect the input ports to the output ports.
3. **Output Ports:** Memory buffers that temporarily store packets received from the switching fabric before transmitting them on the outgoing physical link.
4. **Routing Processor:** Executed in software, it runs routing protocols, maintains network link-state information, and computes the main forwarding table.

3. Detailed Component Breakdown

Input Port Processing & Fast Lookup

The input port performs three tightly decoupled, sequential steps before handing a packet to the switch fabric:

- **Line Termination:** Manages the physical-layer functionality of terminating an incoming physical line at the router.
- **Data Link Processing:** Handles decapsulation and link-layer protocol rules required to communicate with the opposite side of the physical link.
- **Lookup and Forwarding:** Uses the destination IP address to index into a local copy of the forwarding table to pick the proper outgoing port.

Avoiding Central Processing Bottlenecks: Rather than querying the centralized routing processor for every incoming packet, the master forwarding table is continuously copied from the routing processor directly to the line cards over an internal bus (e.g., a PCI bus). This local "shadow copy" allows forwarding decisions to be made purely at the individual line card level.

Hardware vs. Software Timescales

Because a router's input ports, output ports, and switching fabric execute the high-frequency data plane actions, they are almost universally implemented directly in hardware.

- **The Nanosecond Scale (Hardware Data Plane):** Consider a modern 10 Gbps input link processing a standard 64-byte IP datagram. The input port has a window of only **51.2 nanoseconds** to process that packet before the subsequent one arrives. If a line card bundles multiple ports together, the required processing pipeline must execute even faster, rendering software implementations unfeasible.
- **The Millisecond/Second Scale (Software Control Plane):** Control functions—such as running routing protocols (OSPF, BGP), responding to links dropping, and performing remote configurations—operate on much longer timescales. These execute via software residing on a traditional CPU (the Routing Processor).

Would you like to continue to the next topic, which details the structural methods of Switching (via Memory, Bus, and Interconnection Networks) along with Queuing and Head-of-the-Line (HOL) blocking?



Topic 3: Switching Fabrics, Queuing Dynamics, and Packet Scheduling

The switching fabric is the core component of a router that physically transfers packets from an input port to an output port. This section explores the technical mechanisms used to achieve switching, where internal bottlenecking and queuing occur, and how routers schedule packets for transmission.

1. The Three Primary Switching Techniques

Switching within a router can be accomplished using one of three classic architectural designs:

A. Switching via Memory

- **Mechanism:** Early routers were traditional computers where switching operated under the direct control of the central CPU.
- **Operation:** An arriving packet triggers an interrupt to the routing processor. The packet is copied from the input port into processor memory. The CPU extracts the destination address, locates the output port in the forwarding table, and copies the packet to the output port's buffers.
- **Limitations:** Throughput is strictly limited by memory bandwidth. If the memory can handle a maximum of B packet read/write operations per second, the total forwarding throughput must be less than $B/2$. Two packets cannot be transferred simultaneously over the shared system bus.

- **Modern Variation:** Modern memory-based routers shift the lookup and direct memory allocation tasks onto the input line cards themselves, allowing them to function like shared-memory multiprocessors. Example: Cisco's Catalyst 8500 series.

B. Switching via a Bus

- **Mechanism:** Input ports transfer packets directly to output ports over a shared backplane bus without any CPU intervention.
- **Operation:** The input port prepends a switch-internal label (header) to the packet indicating the local output port. The packet is transmitted onto the bus. While all output ports receive the packet, only the port matching the internal label keeps it. The label is then stripped at the output port.
- **Limitations:** If multiple packets arrive simultaneously at different input ports, all but one must wait because only a single packet can cross the bus at any given moment. Forwarding speed is entirely bottlenecked by the physical speed of the single bus. Example: The Cisco 6500 switches packets over a 32-Gbps-backplane bus.

C. Switching via an Interconnection Network

- **Mechanism:** Uses a multi-bus structure to overcome the bandwidth limitations of a single shared bus.
- **Operation:** A crossbar switch consists of an interconnected network of $2N$ buses connecting N input ports to N output ports. A switch fabric controller dynamically opens or closes electronic crosspoints at the intersections of these buses to complete data paths.
- **Advantages:** Capable of forwarding multiple packets in parallel if their target output paths are completely different. Crossbar switches are non-blocking—a packet is never blocked from reaching an output port as long as no other packet is concurrently being forwarded to that same output port.
- **Scaling:** High-performance designs deploy multi-stage switching fabrics (e.g., Cisco CRS uses a three-stage non-blocking approach) or operate multiple fabrics in parallel, breaking packets into smaller chunks and "spraying" them across fabrics to be reassembled at the destination. Example: Cisco 12000 series.

2. Queuing Dynamics: Where Backups Occur

Packet queues form at both input and output ports depending on the relative speed of the switching fabric (R_{switch}), the input line transmission rate (R_{line}), and the current traffic load. When these queues grow large enough to exhaust available memory, packet loss occurs; this is exactly where packets are dropped "within the network".

Input Queuing & Head-of-the-Line (HOL) Blocking

If the switch fabric rate (R_{switch}) is at least N times faster than the line rate (R_{line}), input queuing is negligible. However, if the fabric is not fast enough, arriving packets must queue at the input port. This induces **Head-of-the-Line (HOL) blocking**:

HOL Blocking Definition: A queued packet at the front of an input queue blocks subsequent packets behind it from transferring through the fabric, even if those subsequent packets are destined for free, un-congested output ports.

Under certain traffic assumptions, HOL blocking can cause input queues to grow infinitely long (resulting in heavy packet loss) as soon as the total packet arrival rate hits just 58% of the input link capacity.

Output Queuing & Buffer Management

Even if the switching fabric is N times faster than the line speed, queues will still form at the output ports if multiple packets arriving across different input ports are concurrently destined for the exact same output port.

When output port memory is completely exhausted, the router must choose a discard policy:

- **Drop-Tail:** The router drops the newly arriving packet to protect currently buffered data.
- **Active Queue Management (AQM):** Routers proactively drop or mark packet headers before the buffer fills up completely to send an early congestion signal to sender protocols. **Random Early Detection (RED)** is a widely implemented AQM

algorithm used for this purpose.

Buffer Sizing Formulas

Rule Basis	Buffer Size Formula (B)	Context / Conditions
Traditional Rule of Thumb	$B = \text{RTT} \cdot C$	Derived from analyzing a small number of TCP flows. (e.g., a 10 Gbps link (C) with a 250 ms round-trip time (RTT) requires 2.5 Gbits of buffer memory).
Modern Large-Scale Rule	$B = \frac{\text{RTT} \cdot C}{\sqrt{N}}$	Applicable when a very large number of independent TCP flows (N) pass through backbone router links.

3. Packet Scheduling Disciplines

Once packets are buffered at the output port, a packet scheduler determines the exact transmission order over the outgoing link.

A. First-In, First-Out (FIFO)

- **Behavior:** Selects packets for link transmission in the exact order in which they arrived at the output queue.
- **Mechanics:** Completely ignores packet classes or header field values; once a packet finishes service, the next packet at the front of the line is served.

B. Priority Queuing

- **Behavior:** Packets are separated into distinct priority classes upon arrival.
- **Mechanics:** The scheduler will always transmit packets from the highest-priority class queue that is non-empty. Packets within the same priority class are handled in a FIFO manner.
- **Non-preemptive rule:** Once a packet begins transmission, its service cannot be interrupted by a higher-priority packet that arrives mid-transmission.

C. Round Robin

- **Behavior:** Alternates service sequentially among different packet classes rather than enforcing a strict priority hierarchy.
- **Work-Conserving Principle:** A work-conserving queue discipline never allows the link to remain idle if packets are waiting. If a round-robin cycle checks a specific class queue and finds it empty, it immediately moves to evaluate the next class in sequence.

D. Weighted Fair Queuing (WFQ)

- **Behavior:** A highly sophisticated, generalized form of round-robin scheduling heavily deployed in modern routers.
- **Mechanics:** Arriving packets are classified and placed into per-class queues. Each class i is assigned a specific configuration weight w_i .
- **Throughput Guarantee:** During any period where multiple classes have queued packets, each class i is guaranteed to receive a fraction of total bandwidth equal to: $\frac{w_i}{\sum w_j}$. The denominator sum is taken across all active classes with buffered packets. For a link with a transmission rate R , class i is always guaranteed a processing throughput of at least: $R \cdot \frac{w_i}{\sum w_j}$

Please state **continue** to move to the next topic, which covers the Internet Protocol (IP), the IPv4 datagram format, and IPv4 fragmentation mechanisms.

Topic 4: The Internet Protocol (IP) – IPv4 Datagram Format and Fragmentation

The Internet's network-layer packet is formally referred to as a **datagram**. This section covers the structural layout of the IPv4 datagram and the protocol mechanisms used to handle link-layer size constraints through fragmentation.

1. IPv4 Datagram Format

An IPv4 datagram features a fixed base header of 20 bytes when no options are present. If a datagram encapsulates a standard TCP segment, the combined header overhead is 40 bytes (20 bytes of IP header plus 20 bytes of TCP header) alongside the application-layer message.

The individual fields within an IPv4 datagram header are defined as follows:

- **Version Number (4 bits):** Specifies the IP protocol version of the datagram (e.g., 4 for IPv4). Routers check this field to properly interpret the remainder of the packet.
 - **Header Length (4 bits):** Determines where the actual payload begins. This is necessary because an IPv4 datagram can contain a variable number of options.
 - **Type of Service (TOS) (8 bits):** Allows network administrators to distinguish and prioritize different categories of traffic (e.g., separating real-time voice packets from non-real-time file transfers). Two of these bits are also utilized for Explicit Congestion Notification (ECN).
 - **Datagram Length (16 bits):** The total size of the IP datagram (header plus data) measured in bytes. Its theoretical maximum is 65,535 bytes, though practical limits rarely exceed 1,500 bytes to align with standard Ethernet frames.
 - **Identifier, Flags, Fragmentation Offset:** A set of three header fields specifically dedicated to managing IP fragmentation.
 - **Time-to-Live (TTL) (8 bits):** A safety mechanism designed to ensure datagrams do not circulate infinitely due to routing loops. This counter is decremented by 1 at every router that processes the packet; if it reaches 0, the packet is dropped.
 - **Protocol (8 bits):** Indicates the specific transport-layer protocol to which the extracted payload must be delivered at the final destination host (e.g., a value of 6 specifies TCP, while 17 specifies UDP). It serves as the connective glue binding the network and transport layers together.
 - **Header Checksum (16 bits):** Aids routers in detecting bit errors introduced along physical links. It treats each 2-byte boundary in the header as a distinct number, summing them using 1s complement arithmetic. Because the TTL (and potentially options fields) changes at every hop, this checksum must be recomputed and rewritten at every intermediate router. Corrupted datagrams are typically discarded.
 - **Source and Destination IP Addresses (32 bits each):** Stamped by the originating source host, these fields contain the 32-bit globally unique IP addresses of the initial sender and ultimate destination.
 - **Options:** Extensions that allow the standard IP header to carry additional monitoring or routing configuration parameters. Because they introduce variable length and complicate high-speed hardware processing pipelines, options were completely dropped from the newer IPv6 standard.
 - **Data (Payload):** The core content of the packet, typically carrying a transport-layer segment (TCP or UDP) or alternative operational data such as ICMP messages.
-

2. IPv4 Datagram Fragmentation and Reassembly

Different link-layer protocols use different underlying physical frame configurations. The maximum amount of data that a specific link-layer frame can physically transport is called its **Maximum Transmission Unit (MTU)**. For example, Ethernet links

support an MTU of up to 1,500 bytes, while certain wide-area links restrict the MTU to 576 bytes.

The Fragmentation Process

When a router receives an incoming datagram that is larger than the MTU of its outgoing target link, the router breaks the datagram's payload into two or more smaller fragmented datagrams, encapsulates each into a unique link-layer frame, and transmits them onto the link.

End-System Reassembly Strategy

To preserve high router forwarding performance and keep core network infrastructure simple, the designers of IPv4 dictated that intermediate routers **never** reassemble fragments. Instead, fragment reassembly is handled purely by the final end-system destination host before the data is handed up to the transport layer.

To execute this reassembly, the destination host evaluates three specific header fields:

- 1. **Identification Number:** The initial sending host stamps every outgoing datagram with an identifying sequence value. When a router splits a datagram, every resulting fragment copies this exact identification number, allowing the destination to group corresponding fragments together.
- 2. **Flag Bits:** To signal when the final fragment has arrived, the very last fragment has its flag bit set to 0, whereas all preceding fragments have this flag bit set to 1.
- 3. **Fragmentation Offset:** Specifies exactly where each piece fits sequentially within the payload structure of the original unfragmented IP datagram, allowing proper sequencing and detection of missing packets.

Fragmentation Example

A datagram of 4,000 bytes (composed of a 20-byte IP header and 3,980 bytes of pure data payload) hits a router interface connecting to a link with a strict 1,500-byte MTU limit. The router divides the 3,980 bytes of data into three sequential pieces to fit within the 1,500-byte frame payload maximums (each fragment including its own appended 20-byte IP header):

Fragment	Total Length	Data Size	Flag Bit (More Fragments)	Offset Value
Fragment 1	1500 bytes	1480 bytes	1	0
Fragment 2	1500 bytes	1480 bytes	1	185 (measured in 8-byte blocks: $1480/8$)
Fragment 3	1040 bytes	1020 bytes	0	370 (measured in 8-byte blocks: $(1480 + 1480)/8$)

Please state **continue** to move to the next topic, which focuses on IPv4 Addressing, Subnet Masks, Classless Interdomain Routing (CIDR), and DHCP.



Topic 5: IPv4 Addressing, Subnet Masks, CIDR, and DHCP

1. Fundamental Concepts of IP Addressing

- **Definition of an Interface:** A host typically connects to a network via a single physical link, and the boundary between the host computer and that link is called an **interface**. Routers have the task of receiving datagrams on one link and forwarding them on another, meaning a router necessarily connects to two or more links and possesses multiple corresponding interfaces.
- **Address Association:** Because every host and router interface must be capable of sending and receiving datagrams, IP requires each unique interface to have its own assigned IP address. Therefore, an IP address is technically associated with a specific interface rather than the physical host or router device containing it.
- **Structure and Notation:** Each IPv4 address is exactly 32 bits long (or 4 bytes), which yields a total theoretical pool of 2^{32} (approximately 4 billion) possible unique addresses. To make them human-readable, addresses are typically written in **dotted-decimal notation**, where each individual byte of the address is written in its decimal form (ranging from 0 to 255) and separated from other bytes by a period.

2. IP Subnets and the "Subnet Recipe"

- **Subnet Definition:** In Internet Protocol terms, a network that interconnects multiple host interfaces and a router interface forms a **subnet** (also referred to as an IP network or simply a network). A key characteristic of a subnet is that the underlying network layout connects these interfaces together without any intermediate routers.
- **Subnet Masks:** IP addressing assigns a structure to this subnet using a slash notation (e.g., `/24`), known as a **subnet mask**. The `/24` notation indicates that the leftmost 24 bits of the 32-bit address value strictly define the subnet address boundary. Any additional host interface attached to that same subnet is required to have an address that shares those identical leftmost network bits.
- **The Subnet Recipe:** For any complex interconnected system of routers and hosts, a standard recipe can be used to isolate and define the subnets present:
 1. Imagine detaching each individual interface from its parent host or router device.
 2. This action creates isolated network segments ("islands"), where the detached interfaces terminate the endpoints of those isolated structures.
 3. Each distinct isolated network island constitutes a single **subnet**.

3. Classless Interdomain Routing (CIDR) vs. Classful Addressing

- **CIDR Mechanism:** The global Internet's primary address assignment strategy is known as Classless Interdomain Routing (CIDR). Under CIDR, the 32-bit IP address is divided into a format of `a.b.c.d/x`, where the `x` most significant bits constitute the network portion of the address and are explicitly referred to as the **network prefix**.
- **Route Aggregation:** An organization is typically assigned a block of contiguous addresses that share a common prefix. When an outside router forwards a datagram destined inside that organization, it only evaluates the leading `x` prefix bits. This ability to use a single prefix to advertise multiple internal networks is called **address aggregation** (or route summarization), which vastly minimizes the size of core internet forwarding tables.

Addressing Scheme	Characteristics & Allocation Style	Architectural Drawbacks & Impact
Classful Addressing (Legacy)	Rigidly forced the subnet prefix portion of an IP address to be exactly 8, 16, or 24 bits in length (known as Class A, B, and C networks, respectively).	Proved highly problematic for supporting the rapid growth of mid-sized organizations. A Class C network could only accommodate $2^8 - 2 = 254$ hosts, which was far too small. Conversely, a Class B network supported up to 65,534 interfaces, which was way too large.

Addressing Scheme	Characteristics & Allocation Style	Architectural Drawbacks & Impact
CIDR (Modern)	Allows variable-length network prefixes (/x) tailored directly to an organization's specific scale requirements.	Resolved the rapid depletion of the Class B address space and minimized poor utilization of assigned blocks.

- **The IP Broadcast Address:** A specific, highly critical IP address is 255.255.255.255 . When a host transmits a datagram with its destination field targeted exactly to 255.255.255.255 , the message is automatically delivered to every single host interface residing on that identical local subnet.


4. Dynamic Host Configuration Protocol (DHCP)

While router interfaces are manually configured by network administrators, host interfaces dynamically acquire their address spaces using DHCP.

- **Core Role:** DHCP allows a host to automatically obtain an assigned IP address from a server each time it connects to the network infrastructure. Because it automates these complex networking variables seamlessly, it is widely termed a **plug-and-play** or **zeroconf** (zero-configuration) protocol.
- **Configuration Parameters:** In addition to provisioning a temporary or static host IP address, a DHCP server delivers a suite of essential structural information to the arriving node:
 - The interface's assigned **subnet mask**.
 - The IP address of its first-hop router (commonly known as the **default gateway**).
 - The IP address of its local **DNS server**.

The Four-Step DHCP Exchange Protocol

DHCP executes as a client-server protocol over UDP via specific destination ports. For a newly arriving host, the configuration loop follows a strict four-step transaction handshake:



1. **DHCP Server Discovery:** The client constructs a DHCP discover message inside a UDP packet sent to port 67. Because it does not yet know its own network address or the server's location, it uses a generic source address of 0.0.0.0 and a broadcast destination of 255.255.255.255 . The link layer then broadcasts this frame to all local nodes.
2. **DHCP Server Offer:** Any listening DHCP servers on the subnet respond with a DHCP offer message broadcasted using 255.255.255.255 . This offer carries the discovery message's transaction ID, a proposed client IP address (yiaddr), the matching network mask, and an IP address lease time defining how long the allocation remains valid.
3. **DHCP Request:** The client evaluates the available server offers, selects one, and responds with a DHCP request message, broadcasting its parameters back to echo and confirm its configuration parameters.
4. **DHCP ACK:** The chosen server responds to the request with a final DHCP ACK message, confirming all requested parameters and completing the binding transaction.

A Major Shortcoming of DHCP Mobility: Because a distinct IP address is assigned by DHCP each time a node attaches to a different subnet, a stable, end-to-end TCP connection to a remote application cannot be naturally maintained if a mobile host physically moves across subnets. Moving subnets breaks the active socket mapping, forcing the connection to tear down.

Please state **continue** to move to the next topic, which focuses on Network Address Translation (NAT), Middleboxes, and Firewalls / Intrusion Detection Systems (IDS).



Topic 6: Network Address Translation (NAT) and Middleboxes

1. Motivation for NAT

The proliferation of small office, home office (SOHO) subnets created a massive demand for IP addresses. If every device in every home (such as phones, tablets, gaming devices, and printers) required a unique global IPv4 address, ISPs would constantly need to allocate large blocks of contiguous addresses, which is highly inefficient and quickly depletes the available address space.

Network Address Translation (NAT) provides a simpler, more manageable approach to address allocation for these scenarios.

2. Private Address Realms

- **Private Address Spaces:** Under RFC 1918, specific portions of the IP address space are reserved strictly for private networks. A common example is the 10.0.0.0/8 block (often used as 10.0.0.0/24 in home networks).
 - **Definition of a Private Realm:** A private address realm refers to a network whose internal addresses only have operational meaning to the devices residing within that specific network.
 - **The Isolation Problem:** Hundreds of thousands of home networks use the exact same address space (e.g., 10.0.0.1). Because these addresses are not globally unique, packets leaving a home network cannot use private addresses on the global Internet; doing so would cause routing conflicts across identical network boundaries.
-


3. NAT Operation and the Translation Table

To the outside world, a NAT-enabled router does not look or behave like a standard router. Instead, it acts as a single device with one globally unique public IP address. All traffic leaving the home network for the global Internet uses this single public WAN interface address, effectively hiding the internal architectural details of the home network.

The Translation Mechanism

Because all incoming and outgoing datagrams share a single public IP address, the router uses a **NAT translation table** to map traffic back to the correct internal host. This table includes both IP addresses and transport-layer port numbers in its entries. Since a port number field is 16 bits long, a single public IP address can support over 60,000 simultaneous network connections.

Step-by-Step Packet Walkthrough



- 1. Outgoing Transmission:** A browser on internal host 10.0.0.1 requests a webpage from a web server (128.119.40.186, port 80). The host assigns an arbitrary source port number (e.g., 3345) and sends the datagram into the local LAN.
- 2. NAT Router Mapping:** The NAT router intercepts the datagram. It generates a new source port number (e.g., 5001) that is not currently active in its table, replaces the source IP with its own public WAN IP (138.76.29.7), and replaces the original source port with 5001. It then adds this binding to the translation table and forwards the packet onto the WAN link.
- 3. External Server Reply:** The remote web server receives the packet and responds with a return datagram addressed to the router's public IP (138.76.29.7) and destination port 5001.
- 4. Incoming Translation:** When the response arrives, the NAT router uses the destination IP and port (5001) to index into the translation table. Finding the match for internal host 10.0.0.1 and port 3345, it rewrites the packet headers and delivers the datagram to the correct device inside the home network.

4. Architectural Debates and Architectural Purism

Despite its widespread adoption, NAT faces significant criticism from networking architectural purists:

- **Header Violation:** Critics argue that transport-layer port numbers are meant to address software processes, not identify physical host machines.
- **End-to-End Principle Breakdown:** NAT violates the fundamental networking principle that hosts should communicate directly with one another without intermediate nodes modifying IP addresses or port fields.
- **Inbound Connection Issues:** Because the translation table requires an outgoing packet to establish a mapping, external hosts cannot transparently initiate inbound connections to servers or peers operating behind a NAT. Technical workarounds include **NAT traversal tools** (such as STUN) and **Universal Plug and Play (UPnP)** protocols.

5. Middleboxes and Network Security

NAT routers belong to a broader category of network components called **middleboxes**. Middleboxes operate at the network layer but perform functions other than traditional packet forwarding, including load balancing, firewalling, and deep packet inspection.

Firewalls

Firewalls inspect network layer datagrams and transport segment header fields to deny entry to suspicious packets. They protect internal networks by:

- Blocking packets based on specified source/destination IP addresses or port ranges.
- Dropping ICMP messages (like ping sweeps) to prevent attackers from mapping the network's active IP ranges.
- Tracking the state of active TCP connections, granting entry only to datagrams that explicitly match pre-approved, safely established connections.

Intrusion Detection and Prevention Systems (IDS / IPS)

- **Intrusion Detection System (IDS):** Positioned at network boundaries, an IDS performs **deep packet inspection (DPI)** by examining both header fields and the actual data payload (including application-layer fields). It matches this data against a database of known attack signatures and generates security alerts when a match occurs.
- **Intrusion Prevention System (IPS):** An IPS functions identically to an IDS, but has the added capability to actively block malicious packets from entering the network rather than just raising a passive alert.

You Asked 

continue



Topic 7: The IPv6 Protocol and Transition from IPv4

1. Motivation for IPv6

In the early 1990s, the Internet Engineering Task Force (IETF) began developing a successor to the IPv4 protocol. The prime motivation was the realization that the 32-bit IPv4 address space was rapidly being depleted due to the breathtaking rate at which new subnets and nodes were connecting to the Internet. In February 2011, the Internet Assigned Numbers Authority (IANA) officially allocated the last remaining pool of unassigned IPv4 addresses to regional registries.

Beyond expanding address capacity, the developers of IPv6 used the opportunity to streamline and optimize the protocol based on decades of operational experience with IPv4.

2. Key Architectural Changes in IPv6

- **Expanded Addressing Capabilities:** IPv6 increases the size of an IP address from 32 bits to 128 bits. This massive expansion ensures the world will not run out of addresses, effectively making every grain of sand on the planet addressable.
- **Anycast Addresses:** Alongside standard unicast and multicast formats, IPv6 introduces an *anycast address*. This allows a datagram to be delivered to any single member of a specified group of hosts (e.g., routing an HTTP request to the closest available document mirror site).
- **Streamlined 40-Byte Fixed Header:** Several fields from the traditional IPv4 header were dropped entirely or converted into optional fields. The resulting 40-byte fixed-length header enables significantly faster processing by routers.
- **Flow Labeling:** IPv6 includes a dedicated 20-bit field to identify a specific "flow" of datagrams. This allows a sender to request special handling from routers, such as prioritizing non-default quality of service (QoS) or real-time audio/video transmission over standard data transfers like email.

3. IPv6 Datagram Header Format

The individual header fields of an IPv6 datagram include:

- **Version (4 bits):** Identifies the IP version number; for IPv6, this field holds a value of 6.

- **Traffic Class (8 bits):** Functions similarly to the IPv4 Type of Service (TOS) field, allowing routers to grant priority to specific applications or datagrams within a flow.
 - **Flow Label (20 bits):** Identifies a distinct packet flow requiring specialized handling.
 - **Payload Length (16 bits):** An unsigned integer indicating the total number of bytes in the datagram following the fixed 40-byte header.
 - **Next Header (8 bits):** Identifies the protocol to which the payload will be delivered (e.g., TCP or UDP), utilizing the same protocol numbers as IPv4. It can also point to a standalone options field chained after the main header.
 - **Hop Limit (8 bits):** A safety counter decremented by 1 at each routing hop; the datagram is instantly discarded if the count hits zero.
 - **Source and Destination Addresses (128 bits each):** The globally unique source and destination endpoints of the datagram.
-

4. Fields Removed from the IPv4 Specification

To maximize processing performance, structural operations that burdened intermediate routers were eliminated or moved to the network edge:

- **Fragmentation and Reassembly:** IPv6 completely forbids intermediate routers from performing fragmentation or reassembly; these actions can only occur at the source and destination end-systems. If an incoming IPv6 datagram is too large for an outgoing link, the router drops it and sends an ICMP "Packet Too Big" error message back to the sender, prompting a retransmission with smaller sizes.
 - **Header Checksum:** Because transport-layer (TCP/UDP) and link-layer (Ethernet) protocols already perform comprehensive error checking, the network-layer checksum was dropped as redundant. This saves intermediate routers from having to recompute a checksum at every single hop when the Hop Limit field changes.
 - **Options:** Dropping options from the standard header results in a strict, predictable fixed-length layout. If options are required, they are treated as an independent payload layer indicated by the **Next Header** pointer.
-

5. Transitioning from IPv4 to IPv6: Tunneling

A sudden global transition date (a "flag day") where all network devices are concurrently upgraded is completely impossible given the billions of devices active on the modern Internet. Instead, the most widely adopted operational bridge between the two protocols is **tunneling**.

The Tunneling Mechanism

When two IPv6-capable routers need to communicate but are separated by an intervening network of traditional IPv4 routers, they treat the IPv4 infrastructure as a point-to-point data link layer—a "tunnel".

Step-by-Step Tunneling Operation

1. **Encapsulation:** The sending IPv6 node takes the *entire* native IPv6 datagram and inserts it into the data (payload) field of a standard IPv4 datagram.
 2. **Addressing:** The outer IPv4 header is addressed directly to the public IPv4 interface of the receiving IPv6 node at the opposite end of the tunnel.
 3. **Transit:** Intervening IPv4 routers forward this packet among themselves normally, completely unaware that the inner payload contains an active IPv6 packet.
 4. **Decapsulation:** When the destination router receives the outer IPv4 packet, it observes a protocol number field value of **41**. This specifies that the payload is an encapsulated IPv6 datagram. The router extracts the inner IPv6 datagram and processes it along its native path.
-

Would you like to proceed to the final topic, which covers Generalized Forwarding, the OpenFlow Match-Plus-Action abstraction, and Software-Defined Networking (SDN) examples?



Topic 8: Generalized Forwarding and SDN (The Match-Plus-Action Abstraction)

In traditional destination-based routing, a router matches the destination IP address of an incoming packet against its local forwarding table to pick a specific output interface. Modern network architectures—specifically **Software-Defined Networking (SDN)**—utilize a more powerful concept known as **generalized forwarding**. Here, forwarding decisions can be based on multiple fields across different protocol layers, not just the destination network address.

1. Destination-Based vs. Generalized Forwarding

- **Destination-Based Forwarding:** A packet switch looks up the destination IP address in the header and uses longest-prefix matching to determine the single outbound link interface.
- **Generalized Forwarding:** A packet switch evaluates a diverse set of header fields across the link, network, and transport layers. It matches these headers against entries in a **flow table** to execute a wide variety of behaviors, such as dropping the packet, rewriting header values, or duplicating the packet across multiple ports.

2. The OpenFlow Match-Plus-Action Abstraction

OpenFlow is the pioneering standard protocol that defines how an SDN remote controller communicates with generalized hardware switches. In OpenFlow, a router or switch is conceptualized as a programmable table of rules.

Each entry in a **Flow Table** consists of three primary components:

1. **A Set of Match Fields:** The specific header criteria that an incoming packet must satisfy to trigger the rule.
2. **A Set of Actions:** The operational tasks executed by the switch hardware if a packet matches all criteria in the entry.
3. **A Set of Counters:** Variables that track statistical information, such as the total number of packets or bytes that have successfully matched this specific rule (crucial for traffic profiling and billing metrics).

3. Detailed Breakdown of OpenFlow Elements

A. The Match Fields

OpenFlow allows rules to match fields spanning three distinct layers of the OSI protocol stack:

- **Link Layer (Layer 2):** Source and Destination MAC addresses, Ethernet frame type, and VLAN ID tags.
- **Network Layer (Layer 3):** Source and Destination IPv4/IPv6 addresses, and the IP Protocol field (e.g., TCP vs. UDP).
- **Transport Layer (Layer 4):** TCP or UDP Source and Destination port numbers.

Wildcards: OpenFlow supports wildcard entries. For example, a match field configuration of `IP Src = 192.168.*.*` matches any incoming packet originating from that broad internal subnet, ignoring the host portion.

B. The Actions

When a packet matches a flow entry, the switch can execute one or more of the following actions:

- **Forwarding:** Directs the packet to a specific physical output port, broadcasts it across all ports, or encapsulates and routes it directly to the remote SDN controller for deeper software evaluation.

- **Dropping:** Instantly discards the packet (effectively functioning as a firewall rule).
- **Modify-Field:** Rewrites or alters header fields before transmitting the packet over the outbound link (crucial for executing NAT or shifting VLAN boundaries in hardware).

4. Unifying Multiple Network Devices Under One Abstraction

The power of generalized forwarding is that a single piece of hardware can be dynamically programmed via software to mimic completely different network devices:

Desired Network Device	OpenFlow Match Configuration	Corresponding OpenFlow Action
Traditional Router	Matches the destination IP address (using longest-prefix matching).	Forwards the packet to the specified outbound interface port.
Link-Layer Switch	Matches the destination MAC address.	Forwards the frame to the corresponding physical port mapped to that MAC.
Firewall	Matches specific Layer 4 port numbers (e.g., blocking destination port 22 for SSH) or unwanted source IP addresses.	Drops the packet.
Network Address Translation (NAT)	Matches packets based on internal IP addresses and source port numbers.	Rewrites the source IP and port fields to match the public WAN profile, then forwards the packet.

This concludes the detailed notes for **Chapter 4: The Network Layer — Data Plane**.

If you would like to move on to the next major component—**Chapter 5: The Network Layer — Control Plane** (covering Routing Algorithms like Link-State/OSPF, Distance-Vector/RIP, BGP, and SDN Controller Architecture)—please type **continue**.



Topic 9: Problem-Solving Frameworks and Analytical Review

This section synthesizes the core analytical, mathematical, and design principles required to solve the architectural and quantitative problems found at the conclusion of Chapter 4.

1. Longest Prefix Match (LPM) Evaluation Framework

When a destination address matches multiple entries in a routing table, a packet switch must apply the **longest prefix matching rule**. This selects the entry with the longest matching bit prefix.

Step-by-Step Conversion and Matching Technique

To solve problem types like **P5**, **P6**, or **P7**:

1. **Convert to Binary:** Translate both the incoming destination IP address and the routing prefixes into binary strings.
2. **Align and Compare:** Compare bits from left to right.

- Count Matching Bits:** Determine the exact number of matching leading bits for each entry.
- Select Maximum:** Choose the interface corresponding to the entry with the absolute highest count of matching bits. If no explicitly listed prefix matches, default to the wildcard/otherwise path.

```
Incoming Address: 11100001 01000000 11000011 00111100

Table Prefix A: 11100000 * -> 7 matching bits
Table Prefix B: 11100001 01000000 * -> 16 matching bits
Table Prefix C: 11100001 01000001 * -> 15 matching bits (16th bit mismatches)

Selection: Prefix B (Longest Match = 16 bits)
```

2. Subnetting Strategy and Address Allocation Math

Designing subnets to support a specific number of interfaces (such as in **P8** or **P12**) requires calculating host capacity based on prefix lengths ($/x$).

The Host Capacity Formula

For a given subnet block with prefix length x , the number of available bits for host interfaces is $32 - x$. The total number of addressable interfaces is: $\text{Capacity} = 2^{32-x} - 2$ (*Subtracting 2 accounts for the reserved network address where host bits are all 0, and the subnet broadcast address where host bits are all 1*).

Prefix Sizing Rules

Required Host Interfaces	Minimum Address Pool Size	Host Bits Needed ($32 - x$)	Prefix Mask ($/x$)
≤ 2 (Point-to-point)	4	2	/30
≤ 14	16	4	/28
≤ 30	32	5	/27
≤ 62	64	6	/26
≤ 126	128	7	/25
≤ 254	256	8	/24

Design Tip (Variable Length Subnet Masking - VLSM): Always allocate the largest subnet first to prevent overlapping address blocks. For example, if a network requires blocks of 90, 60, and 12 hosts, allocate the **/25** block first, followed by the **/26** block, and finally the **/28** block.

3. IP Fragmentation Calculus

When an IP datagram crosses a link with a smaller Maximum Transmission Unit (MTU), it must be divided into fragments.

Key Math Constraints

- Header Overhead:** Each fragment requires its own 20-byte IP header.
- Maximum Payload Capacity:** $\text{Max Payload} = \text{MTU} - 20$.
- The 8-Byte Alignment Rule:** The payload size of every fragment except the last one **must be a multiple of 8 bytes** to accommodate the 13-bit fragmentation offset field. If Max Payload is not divisible by 8, round *down* to the nearest multiple of 8.
- Offset Value Calculation:** $\text{Offset} = \frac{\text{Total accumulated payload bytes sent in previous fragments}}{8}$

Example Walkthrough (2400-byte Datagram into a 700-byte MTU Link)

- **Original Packet:** 20 bytes header + 2380 bytes payload.
- **Link Constraint:** $MTU = 700 \implies \text{Max raw payload} = 700 - 20 = 680 \text{ bytes}$.
- **Check Alignment:** $680/8 = 85$ (Perfectly divisible by 8). Each fragment can carry exactly 680 bytes of data.

```
Original Payload: 2380 bytes
|
+---> Fragment 1: Data = 680 bytes | Offset = 0    | More Flags = 1 | Total Length = 700
|
+---> Fragment 2: Data = 680 bytes | Offset = 85   | More Flags = 1 | Total Length = 700
|
+---> Fragment 3: Data = 680 bytes | Offset = 170  | More Flags = 1 | Total Length = 700
|
+---> Fragment 4: Data = 340 bytes | Offset = 255 | More Flags = 0 | Total Length = 360
```

4. OpenFlow Match-Plus-Action Flow Table Composition

SDN flow tables require specific configurations to implement custom pathing, load balancing, or network-layer firewalls.

Logical Entry Layout

To design tables like those in **P19** or **P22**, map the criteria systematically using matches and actions:

FLOW TABLE	
MATCH	ACTION
• Ingress Port • Layer 2 (Src/Dst MAC, EthType) • Layer 3 (Src/Dst IP, Protocol) • Layer 4 (Src/Dst Port)	• Forward(Port Number) • Drop • Modify-Field (Rewrite Header)

Multi-Criteria Firewall Example

If a policy requires that **only UDP traffic from Host 1 (10.1.0.1) destined to Host 3 (10.2.0.3) is allowed**, while all other traffic is blocked:

Entry Priority	Match Fields	Action
Priority 2 (High)	IP Src = 10.1.0.1 , IP Dst = 10.2.0.3 , IP Proto = 17 (UDP)	Forward(Output_Port)
Priority 1 (Low)	* (Wildcard match for all remaining traffic)	Drop

This completes the review and problem-solving frameworks for **Chapter 4: The Network Layer — Data Plane**. You are now fully prepared to examine or complete the conceptual study of this chapter. If you have any other files to analyze or wish to explore a specific problem in further detail, please let me know!



Topic 10: Step-by-Step Quantitative Solutions to Core Chapter Problems

This topic compiles detailed analytical walkthroughs for the primary mathematical and structural problems presented at the end of Chapter 4. Mastering these workflows bridges the gap between theoretical data plane architecture and practical network engineering.

1. Crossbar Fabric Slotted Scheduling Optimization (Problem P4)

Problem Context

A non-blocking crossbar switch fabric operates in a slotted, synchronous manner. In a single time slot, at most one datagram can be transferred to a given output port, but different output ports can receive datagrams from distinct input ports simultaneously.

Consider the following head-of-line queue states where input ports contain packets destined for output ports X , Y , and Z :

- **Input Queue 1:** Packets waiting for $[X]$
- **Input Queue 2:** Packets waiting for $[X, Y]$
- **Input Queue 3:** Packets waiting for $[Z, Y]$

A. Finding the Minimal Number of Time Slots

To achieve the absolute minimum time slots, we schedule packets to maximize parallel transmission across the fabric while avoiding output contention:

- **Slot 1:**
 - Transfer the packet for X from Input 1.
 - Transfer the packet for Y from Input 2.
 - Transfer the packet for Z from Input 3.
 - *Result:* Input 1 is cleared, Input 2 has X remaining, and Input 3 has Y remaining.
- **Slot 2:**
 - Transfer the remaining packet for X from Input 2.
 - Transfer the remaining packet for Y from Input 3.
- **Conclusion:** The minimal number of slots required is **2**.

B. Finding the Worst-Case Number of Slots (Without Idle Queues)

In the worst-case scheduling order where output contention is maximized at every step:

- **Slot 1:** Input 1 transfers to X ; Input 2 is blocked waiting for X ; Input 3 transfers to Z .
- **Slot 2:** Input 2 transfers to X ; Input 3 is blocked waiting for Y .
- **Slot 3:** Input 2 transfers to Y ; Input 3 is blocked waiting for Y .
- **Slot 4:** Input 3 transfers to Y .
- **Conclusion:** The maximum number of slots under a sub-optimal schedule is **4**.

2. Prefix Range and Address Pool Computations (Problems P6 & P7)

Walkthrough for an 8-bit Host Address Table (Problem P6)

Consider a router with an 8-bit address space utilizing the following longest prefix matching rules:

Prefix Match	Outbound Interface
00	Interface 0
010	Interface 1
011	Interface 2
10	Interface 2
11	Interface 3

To find the exact address range and size for each interface, expand the prefix to fill all 8 bits:

- **Interface 0 (00_____):**
 - Binary Range: 00000000 to 00111111
 - Decimal Range: **0 to 63**
 - Total Addresses: $2^{(8-2)} = 2^6 = \mathbf{64 \text{ addresses}}$
- **Interface 1 (010_____):**
 - Binary Range: 01000000 to 01011111
 - Decimal Range: **64 to 95**
 - Total Addresses: $2^{(8-3)} = 2^5 = \mathbf{32 \text{ addresses}}$
- **Interface 2 (Combines 011_____ and 10_____):**
 - Segment A (011_____): 01100000 to 01111111 ($\Rightarrow$ 96 to 127)
 - Segment B (10_____): 10000000 to 10111111 ($\Rightarrow$ 128 to 191)
 - Combined Range: **96 to 191**
 - Total Addresses: $32 + 64 = \mathbf{96 \text{ addresses}}$
- **Interface 3 (11_____):**
 - Binary Range: 11000000 to 11111111
 - Decimal Range: **192 to 255**
 - Total Addresses: $2^6 = \mathbf{64 \text{ addresses}}$

3. Advanced Datagram Fragmentation Breakdown (Problem P14)

Problem Specification

A source transmits a **2,400-byte IP datagram** into a network link with an **MTU limit of 700 bytes**. The original datagram is stamped with an identification number of **422**.

Step-by-Step Calculus

1. Separate Header and Payload:

- Original Header Size = 20 bytes.
- Original Payload Data Size = $2400 - 20 = 2380$ bytes.

2. Determine Maximum Fragment Payload:

- Link MTU limit = 700 bytes.

- Maximum data size per fragment = $700 - 20 = 680$ bytes.

3. Verify 8-Byte Alignment Rule:

- $680/8 = 85$. Because 680 is perfectly divisible by 8, each intermediate fragment will carry exactly 680 bytes of data.

4. Calculate Total Fragments:

- Number of fragments = $\lceil 2380/680 \rceil = \lceil 3.5 \rceil = 4$ fragments.

Fragment Header Parameter Field Mapping

Every fragment retains the original datagram Identification value (**422**). The remaining control parameters evaluate as follows:

Fragment Index	Total Length (Bytes)	Data Payload (Bytes)	"More Fragments" Flag	Offset Header Value (Bytes / 8)
Fragment 1	700	680	1 (True)	$0/8 = 0$
Fragment 2	700	680	1 (True)	$680/8 = 85$
Fragment 3	700	680	1 (True)	$(680 \times 2)/8 = 170$
Fragment 4	460	$2380 - (680 \times 3) = 440$	0 (False)	$(680 \times 3)/8 = 255$

4. Protocol Layer Overhead Analysis (Problem P25)

Problem Specification

A real-time network application generates a data chunk of **40 bytes every 20 milliseconds**. Each discrete chunk is sequentially encapsulated into a standard TCP segment and then wrapped within an IPv4 datagram.

Overhead Calculus

1. Identify Protocol Overhead Layers:

- Standard IPv4 Base Header = 20 bytes.
- Standard TCP Base Header = 20 bytes.
- Total Protocol Overhead = $20 + 20 = 40$ bytes.

2. Determine Total Datagram Size:

- Total Size = Overhead + Application Payload
- Total Size = 40 bytes + 40 bytes = **80 bytes**

3. Percentage Breakdown:

- Overhead Percentage = $\left(\frac{40 \text{ bytes overhead}}{80 \text{ bytes total}} \right) \times 100 = 50\%$
- Application Data Percentage = $\left(\frac{40 \text{ bytes application data}}{80 \text{ bytes total}} \right) \times 100 = 50\%$

This completes our systematic notes for **Chapter 4: The Network Layer – Data Plane**.

To begin generating notes for **Chapter 5: The Network Layer – Control Plane** (including Routing Algorithms, OSPF, BGP, SDN Controllers, and ICMP blocks), please say **continue**.

Deep Dive 1: Variable-Length Subnet Masking (VLSM) & Routing Optimization Matrix (Problem P12)

This deep dive focuses on solving a complex network engineering challenge: allocating a single classless address block across multiple subnets with varied interface scales and constructing their corresponding longest-prefix matching (LPM) hardware forwarding tables.

1. The Design Challenge & Constraints

We are given a single master network prefix: `214.97.254.0/23`. A `/23` mask leaves $32 - 23 = 9$ bits for host identification, providing a total pool of $2^9 = 512$ physical IP addresses. We must distribute this pool across 6 subnets under strict capacity boundaries:

- **Subnet A:** Must support at least **250** host interfaces.
- **Subnet B:** Must support at least **120** host interfaces.
- **Subnet C:** Must support at least **120** host interfaces.
- **Subnets D, E, and F:** Point-to-point router links; each must support exactly **2** interfaces.

The Mathematical Optimization Paradox

If we assign standard classless boundary blocks to match the constraints:

- Subnet A needs 250 hosts → next power of two is 256 → Requires a `/24` block.
- Subnet B needs 120 hosts → next power of two is 128 → Requires a `/25` block.
- Subnet C needs 120 hosts → next power of two is 128 → Requires a `/25` block.

Summing these blocks: 256 (Subnet A) + 128 (Subnet B) + 128 (Subnet C) = 512 addresses. This completely consumes the `/23` address pool, leaving zero addresses for the point-to-point subnets D, E, and F.

The Engineering Solution: Block Carving

To overcome this limitation, we exploit the unused addresses within the allocated power-of-two boundaries. Since Subnet A only requires 250 interfaces out of 254 usable options, a small 4-address block (`/30`) can be safely "carved out" from its upper edge without violating the 250-interface constraint. This technique is repeated across the other subnets to clear space for the link subnets.

2. Final VLSM Address Allocations


Subnet	Target Allocation Style	Network Address / Prefix Mask	Usable Host Range
Subnet A	Main Host LAN	<code>214.97.254.0/24</code>	<code>214.97.254.1</code> to <code>214.97.254.247</code>
Subnet B	Main Host LAN	<code>214.97.255.0/25</code>	<code>214.97.255.1</code> to <code>214.97.255.123</code>
Subnet C	Main Host LAN	<code>214.97.255.128/25</code>	<code>214.97.255.129</code> to <code>214.97.255.247</code>
Subnet D	Carved from Subnet A	<code>214.97.254.248/30</code>	<code>214.97.254.249</code> to <code>214.97.254.250</code>
Subnet E	Carved from Subnet B	<code>214.97.255.124/30</code>	<code>214.97.255.125</code> to <code>214.97.255.126</code>
Subnet F	Carved from Subnet C	<code>214.97.255.248/30</code>	<code>214.97.255.249</code> to <code>214.97.255.250</code>

Proof of Validity

- **Subnet A Usable Capacity:** $247 - 1 + 1 = 247$ hosts. When combined with its 3 local router interfaces on links D, E, and F, it satisfies the ≥ 250 threshold.
- **Subnets B & C Usable Capacity:** Each provides 123 addresses directly, successfully exceeding the 120 interface minimum requirements.

3. Hardware Forwarding Matrix Layout (LPM)

Based on the physical layout specified in Figure 4.20, three internal routers (R1, R2, and R3) interconnect these islands. Each device maintains local interface ports mapped to neighboring nodes:



Router R1 Forwarding Table

- **Port 1:** Directly attached to Subnet A.
- **Port 2:** Attached to Subnet D (leads to R2 and Subnet B).
- **Port 3:** Attached to Subnet F (leads to R3 and Subnet C).

Destination Prefix	Matching Criteria Length	Outbound Target Interface
214.97.254.0/24	24-bit Network Prefix Match	Interface 1
214.97.255.0/25	25-bit Network Prefix Match	Interface 2 (via Link D)
214.97.255.128/25	25-bit Network Prefix Match	Interface 3 (via Link F)
214.97.254.248/30	30-bit Link Prefix Match	Interface 2
214.97.255.248/30	30-bit Link Prefix Match	Interface 3

Router R2 Forwarding Table

- **Port 1:** Attached to Subnet D (leads to R1 and Subnet A).
- **Port 2:** Attached to Subnet E (leads to R3 and Subnet C).
- **Port 3:** Directly attached to Subnet B.

Destination Prefix	Matching Criteria Length	Outbound Target Interface
214.97.255.0/25	25-bit Network Prefix Match	Interface 3

Destination Prefix	Matching Criteria Length	Outbound Target Interface
214.97.254.0/24	24-bit Network Prefix Match	Interface 1 (via Link D)
214.97.255.128/25	25-bit Network Prefix Match	Interface 2 (via Link E)

Router R3 Forwarding Table

- **Port 1:** Attached to Subnet F (leads to R1 and Subnet A).
- **Port 2:** Attached to Subnet E (leads to R2 and Subnet B).
- **Port 3:** Directly attached to Subnet C.

Destination Prefix	Matching Criteria Length	Outbound Target Interface
214.97.255.128/25	25-bit Network Prefix Match	Interface 3
214.97.254.0/24	24-bit Network Prefix Match	Interface 1 (via Link F)
214.97.255.0/25	25-bit Network Prefix Match	Interface 2 (via Link E)

Deep Dive 2: Software-Defined Network Flow Orchestration (Problem P19)

This deep dive transitions from traditional destination-based lookup tables to programmable Software-Defined Networking (SDN) configurations. We construct an OpenFlow entry matrix for packet switch **s2** in Figure 4.30 to execute complex routing, local delivery, and firewall policies within a unified framework.

1. Interface and Node Topology Blueprint

Based on the network rules defined in Section 4.4.3, the physical and logical links terminating at packet switch **s2** map to the following interfaces:

- **Port 1:** Inbound/Outbound link interface routing towards switch **s3** (the gateway for hosts **h5** and **h6**).
- **Port 2:** Inbound/Outbound link interface routing towards switch **s1** (the gateway for hosts **h1** and **h2**).
- **Port 3:** Edge link directly terminating at host **h3** (IP: 10.2.0.3).
- **Port 4:** Edge link directly terminating at host **h4** (IP: 10.2.0.4).

2. Operational Rule Specification

The network administrator specifies four explicit behaviors for traffic passing through switch **s2** :

1. **Transit Rule 1:** Any datagram entering port 1 from the **h5/h6** subnet (10.3.*.*) destined for the **h1/h2** subnet (10.1.*.*) must be forwarded out port 2.
2. **Transit Rule 2:** Any datagram entering port 2 from the **h1/h2** subnet (10.1.*.*) destined for the **h5/h6** subnet (10.3.*.*) must be forwarded out port 1.
3. **Local Inbound Delivery:** Any datagram arriving on ports 1 or 2 targeted directly to local hosts **h3** or **h4** must be immediately delivered to their respective dedicated ports.
4. **Local Peer Communication:** Hosts **h3** and **h4** must be permitted to transmit packets directly to each other without network restrictions.

3. Programmable OpenFlow Flow Table Layout

By combining multi-layer header evaluation with targeted actions, the complete rule set inside s2 is structured as follows:

Rule Priority	Ingress Port	Source IP Match	Destination IP Match	Triggered Hardware Action
Priority 4	Port 1	10.3.*.*	10.1.*.*	Forward(2) (Route to switch s1)
Priority 4	Port 2	10.1.*.*	10.3.*.*	Forward(1) (Route to switch s3)
Priority 3	Port 1	* (Wildcard)	10.2.0.3	Forward(3) (Deliver directly to host h3)
Priority 3	Port 1	* (Wildcard)	10.2.0.4	Forward(4) (Deliver directly to host h4)
Priority 3	Port 2	* (Wildcard)	10.2.0.3	Forward(3) (Deliver directly to host h3)
Priority 3	Port 2	* (Wildcard)	10.2.0.4	Forward(4) (Deliver directly to host h4)
Priority 2	Port 3	10.2.0.3	10.2.0.4	Forward(4) (Local inter-host path)
Priority 2	Port 4	10.2.0.4	10.2.0.3	Forward(3) (Local inter-host path)
Priority 1	*	*	*	Drop (Default security drop policy)

These deep dives provide clear, step-by-step solutions to the most challenging analytical and programmatic problems in the data plane chapter.

If you are ready to move on to the Chapter 5 Control Plane notes, simply reply with continue.